

# Homework1

October 11, 2018

```
In [168]: import numpy as np
          import pandas as pd
```

## 0.1 Problem 1

## 0.2 Part 1, 2

## 0.3 Part 3

```
In [3]: # Define softmax and ReLu
        relu = np.vectorize(lambda z: np.fmax(0,z))
        softmax = lambda x: np.exp(x)/(np.exp(x).sum(axis=0, keepdims=True))
        def predict(y):
            return np.argmax(y,axis=0)

In [4]: def ff_nn_2_ReLu(x,W1,W2,W3,b1,b2,b3):
        a_1 = np.dot(W1,x) + b1
        h1 = relu(a_1)
        a_2 = np.dot(W2,h1) + b2
        h2 = relu(a_2)
        h2
        a_3 = np.dot(W3,h2) + b3
        y = softmax(a_3)
        return y
```

## 0.4 Part 4

```
In [133]: x = np.array([[1,0,0],[-1,-1,1]])
          W1 = np.array([[1,0],[-1,0],[0,.5]])
          W2 = np.array([[1,0,0],[-1,-1,0]])
          W3 = np.array([[1,1],[0,0],[-1,-1]])
          b1 = np.array([[0],[0],[1]])
          b2 = np.array([[1],[-1]])
          b3 = np.array([[1],[0],[0]])
          y = ff_nn_2_ReLu(x,W1,W2,W3,b1,b2,b3)
          y

Out[133]: array([[0.94649912, 0.84379473, 0.84379473],
                 [0.04712342, 0.1141952 , 0.1141952 ],
                 [0.00637746, 0.04201007, 0.04201007]])
```

For each  $x_i$ , the probability of  $x_i$  in class  $j$  is  $(i, j)^{th}$  term of the matrix. For  $1 \leq i \leq n, 1 \leq j \leq m$ , where  $n$  is number of samples (number of columns in  $x$ ) and  $m$  is number of classes (assuming index starts at 1). Hence, for this input, all  $x_1, x_2, x_3$  will be classified in class 1 (assuming index starts at 0)

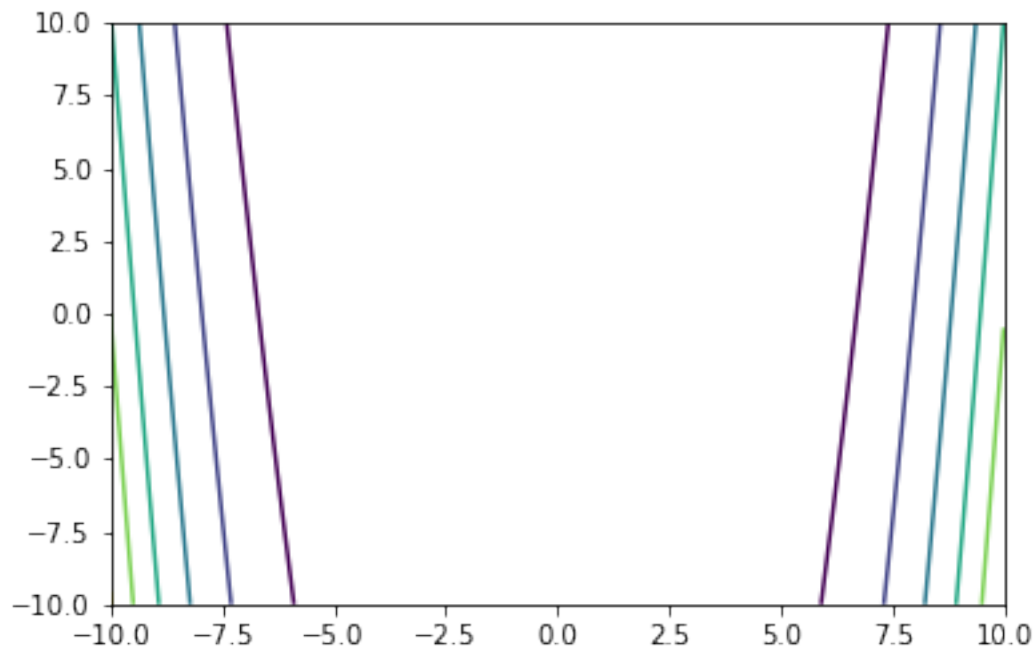
## 0.5 Problem 2

### 0.6 Part 1

$$\frac{\partial f(x, y)}{\partial x} = -2(1 - x) + 200(y - x^2)(-2x) = 2x - 2 - 400x(y - x^2) = 2x - 2 - 400(xy - x^3) \quad \frac{\partial f(x, y)}{\partial y} = 2 * 100(y - x^2) = 200(y - x^2)$$

### 0.7 Part 2

```
In [19]: import matplotlib.pyplot as plt
import numpy as np
# --- Paraboloid Function ---
delta = 0.025
x = np.arange(-10.0, 10.0, delta)
y = np.arange(-10.0, 10.0, delta)
X, Y = np.meshgrid(x, y)
Z = (1-X)**2 + 100*(Y-X**2)**2
fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
```



```
In [101]: # --- Setting up gradient ---
def grad_f(vector):
```

```

x, y = vector
df_dx = 2*x - 2 - 400*(x*y-x**3)
df_dy = 200*(y-x**2)
return np.array([df_dx, df_dy])

# --- Grad Descent ---
def grad_descent(starting_point=None, iterations=10, learning_rate=6):
    if starting_point:
        point = starting_point
    else:
        point = np.random.uniform(-5,5,size=2)
        print('Initial point is',point)
    trajectory = [point]

    for i in range(iterations):
        grad = grad_f(point)
        point = point - learning_rate * grad
        trajectory.append(point)
    print('After iterations, the point is ',point)
    return np.array(trajectory)

```

```

In [102]: # --- Visualize Trajectory ---
np.random.seed(10)
traj = grad_descent(iterations=50,learning_rate=10)

fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
x= traj[:,0]
y= traj[:,1]
plt.plot(x,y,'-o')

```

```

Initial point is [ 2.71320643 -4.79248051]
After iterations, the point is [nan nan]

```

```

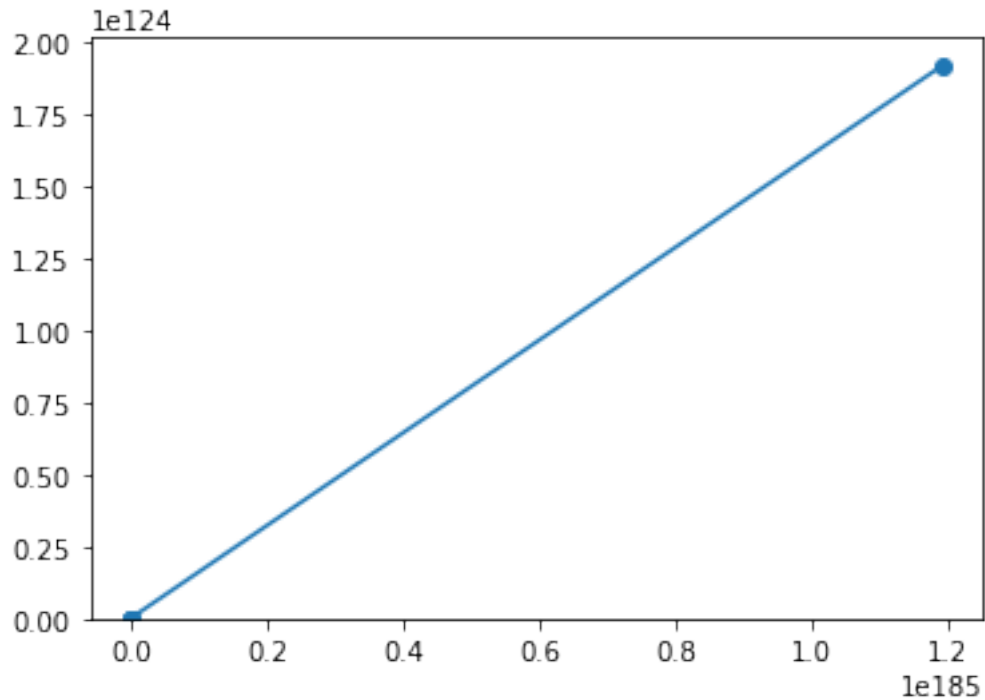
/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:4: RuntimeWarning
  after removing the cwd from sys.path.
/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:4: RuntimeWarning
  after removing the cwd from sys.path.
/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:5: RuntimeWarning
  """

```

```

Out[102]: [<matplotlib.lines.Line2D at 0x144754978>]

```



```
In [103]: # --- Visualize Trajectory ---
          np.random.seed(10)
          traj = grad_descent(iterations=50, learning_rate=1)

          fig, ax = plt.subplots()
          CS = ax.contour(X, Y, Z)
          x= traj[:,0]
          y= traj[:,1]
          plt.plot(x,y, '-o')
```

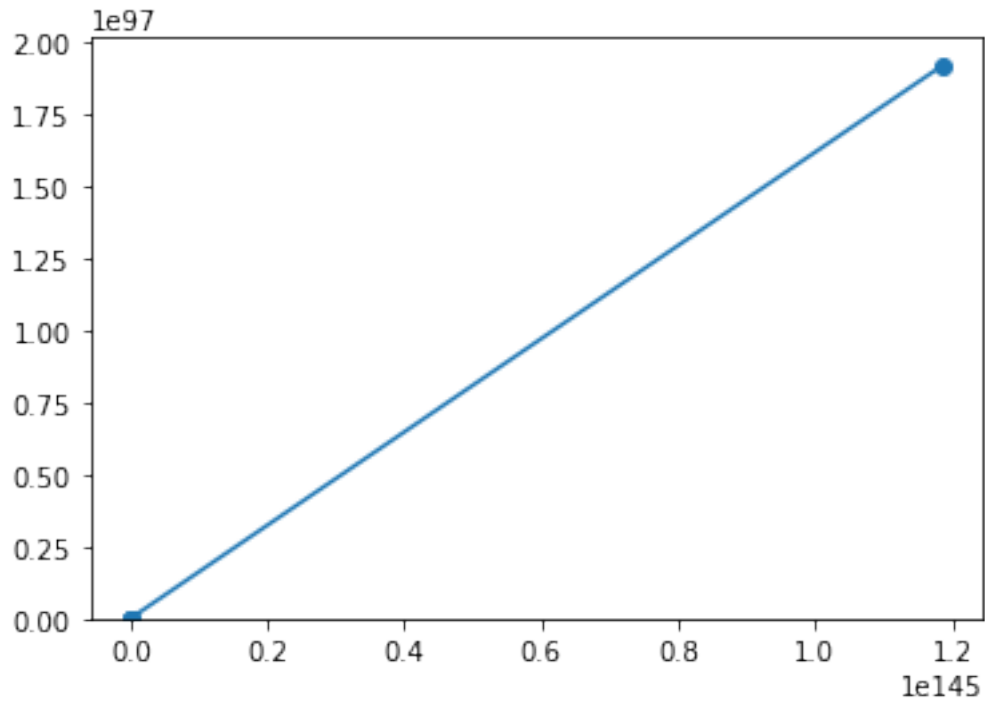
Initial point is [ 2.71320643 -4.79248051]

After iterations, the point is [nan nan]

/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel\_launcher.py:4: RuntimeWarning  
after removing the cwd from sys.path.

/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel\_launcher.py:4: RuntimeWarning  
after removing the cwd from sys.path.

```
Out[103]: [<matplotlib.lines.Line2D at 0x144816a90>]
```



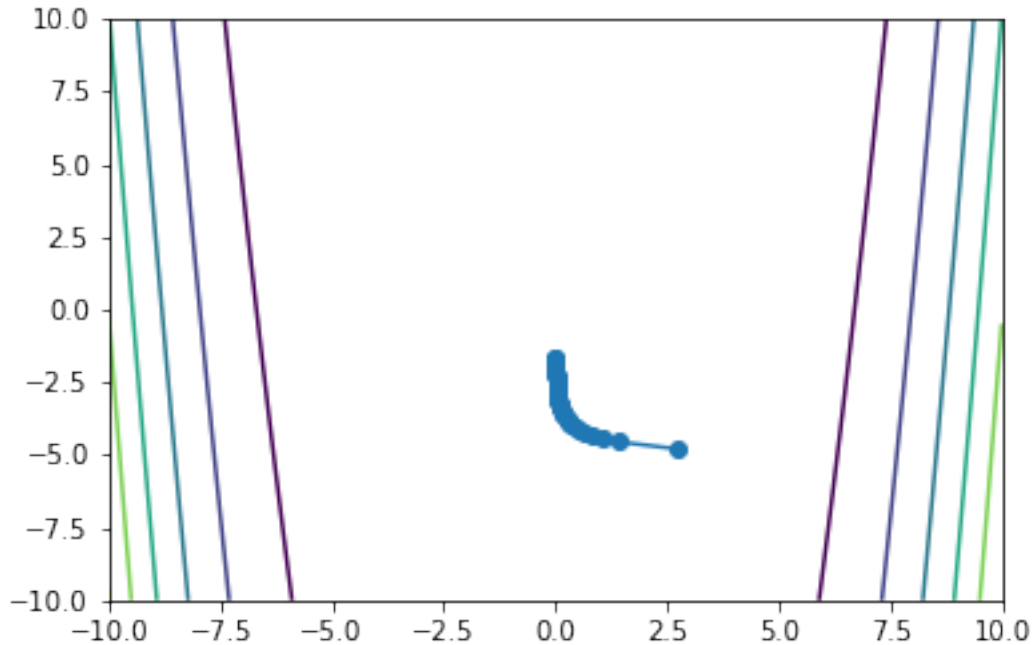
```
In [104]: # --- Visualize Trajectory ---
          np.random.seed(10)
          traj = grad_descent(iterations=50, learning_rate=0.0001)

          fig, ax = plt.subplots()
          CS = ax.contour(X, Y, Z)
          x= traj[:,0]
          y= traj[:,1]
          plt.plot(x,y, '-o')
```

Initial point is [ 2.71320643 -4.79248051]

After iterations, the point is [ 0.00513221 -1.65101579]

```
Out[104]: [<matplotlib.lines.Line2D at 0x1448fba20>]
```



This function is not convex and the minimum point is at  $(-1, -1)$ . When the learning rate is large, the trajectories diverge. Large learning rates will skip the optimal point. On the other hand, when learning rate is small, it becomes more stable (this can be seen from the last trajectory). Here, small learning rate takes each step smaller and it will not skip the optimal point, but it requires more iterations to find the optimal point. Hence, choosing a right learning rate (tuning parameter) is important

## 0.8 Part 4

## 0.9 Gradient Descent with Momentum algorithm

```
In [105]: # --- Gradient with momentum ---
def grad_descent_with_momentum(starting_point=None, iterations=10, alpha=.9, epsilon=1e-6):
    if starting_point:
        point = starting_point
    else:
        point = np.random.uniform(-5,5,size=2)
    trajectory = [point]
    v = np.zeros(point.size)
    print('Initial point is',point)

    for i in range(iterations):
        grad = grad_f(point)
        v = alpha*v + epsilon*grad
        point = point - v
        trajectory.append(point)
```

```

print('After iterations, the point is ',point)
return np.array(trajecory)

```

```

In [106]: # --- Visualizing trajectory ---
np.random.seed(10)
traj = grad_descent_with_momentum(iterations=100, epsilon=2, alpha=.7)

fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
x= traj[:,0]
y= traj[:,1]
plt.plot(x,y,'-o')

```

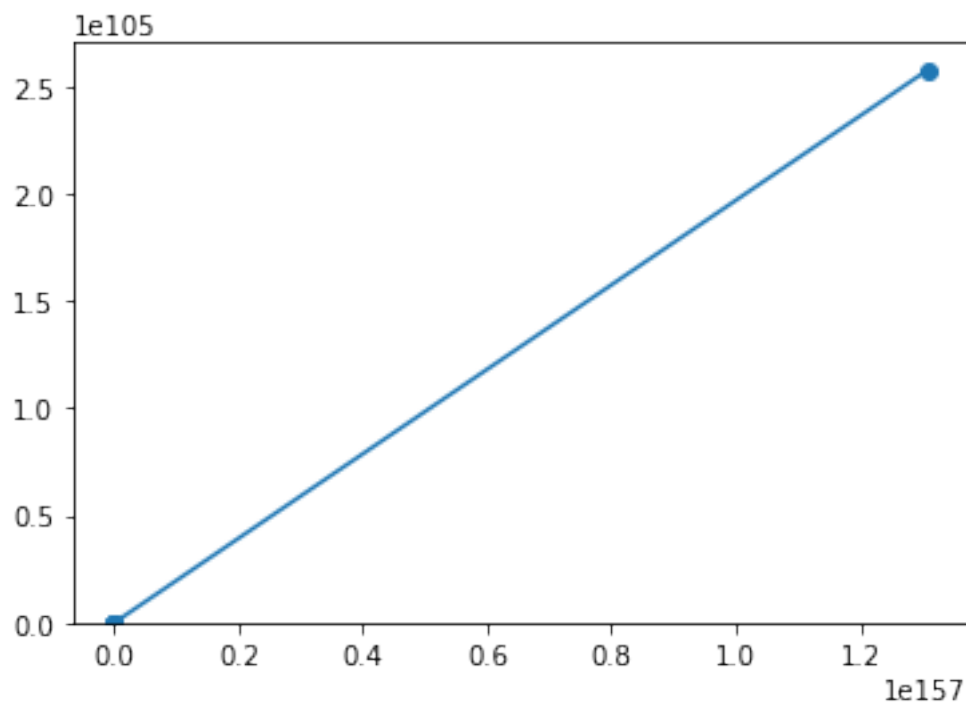
Initial point is [ 2.71320643 -4.79248051]  
After iterations, the point is [nan nan]

```

/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:4: RuntimeWarning
after removing the cwd from sys.path.
/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:5: RuntimeWarning
"""
/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:4: RuntimeWarning
after removing the cwd from sys.path.
/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:5: RuntimeWarning
"""

```

Out[106]: [<matplotlib.lines.Line2D at 0x144816518>]



```

In [107]: # --- Visualizing trajectory ---
          np.random.seed(10)
          traj = grad_descent_with_momentum(iterations=50, epsilon=1, alpha=.5)

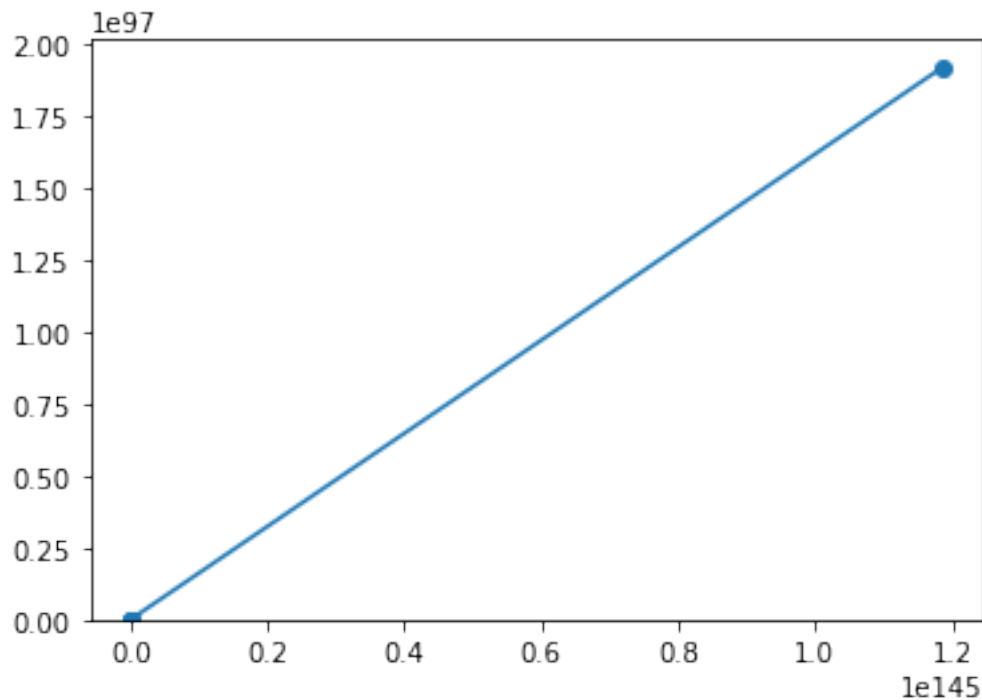
          fig, ax = plt.subplots()
          CS = ax.contour(X, Y, Z)
          x= traj[:,0]
          y= traj[:,1]
          plt.plot(x,y,'-o')

```

Initial point is [ 2.71320643 -4.79248051]  
 After iterations, the point is [nan nan]

/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel\_launcher.py:4: RuntimeWarning  
 after removing the cwd from sys.path.  
 /Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel\_launcher.py:4: RuntimeWarning  
 after removing the cwd from sys.path.

Out[107]: [<matplotlib.lines.Line2D at 0x144a820b8>]



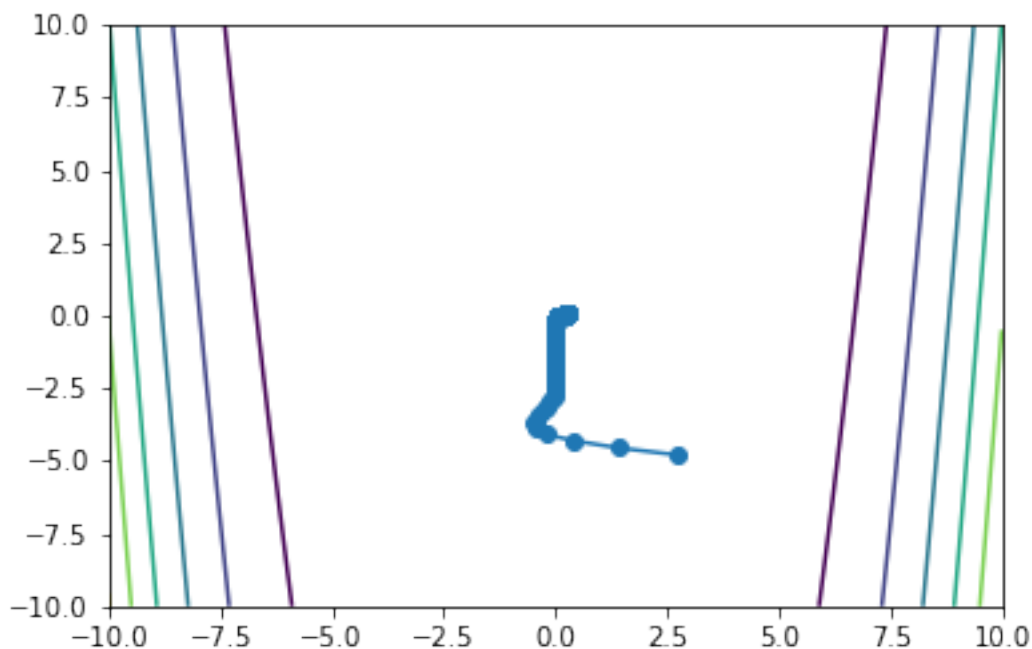


```
In [118]: # --- Visualizing trajectory --
np.random.seed(10)
traj = grad_descent_with_momentum(iterations=1000, epsilon=0.0001, alpha=.5)

fig, ax = plt.subplots()
CS = ax.contour(X, Y, Z)
x= traj[:,0]
y= traj[:,1]
plt.plot(x,y,'-o')
```

Initial point is [ 2.71320643 -4.79248051]  
After iterations, the point is [0.28519766 0.07826446]

Out[118]: [<matplotlib.lines.Line2D at 0x145215c18>]



From all trajectories above, we can see epsilon is similar to learning rate in part c. Larger epsilon makes the algorithm skip the optimal point and diverge. On the other hand, small epsilon makes it more stable but takes more iterations to reach optimal point. For example, in the last plot, initial point is (2.71,-4.79) and it moves toward optimal point. After 1000 iterations, it ends at (0.285,0.0783). Next, I will try a smaller epsilon and see how it performs.

```
In [119]: # --- Visualizing trajectory --
np.random.seed(10)
traj = grad_descent_with_momentum(iterations=1000, epsilon=0.00001, alpha=.5)

fig, ax = plt.subplots()
```

```

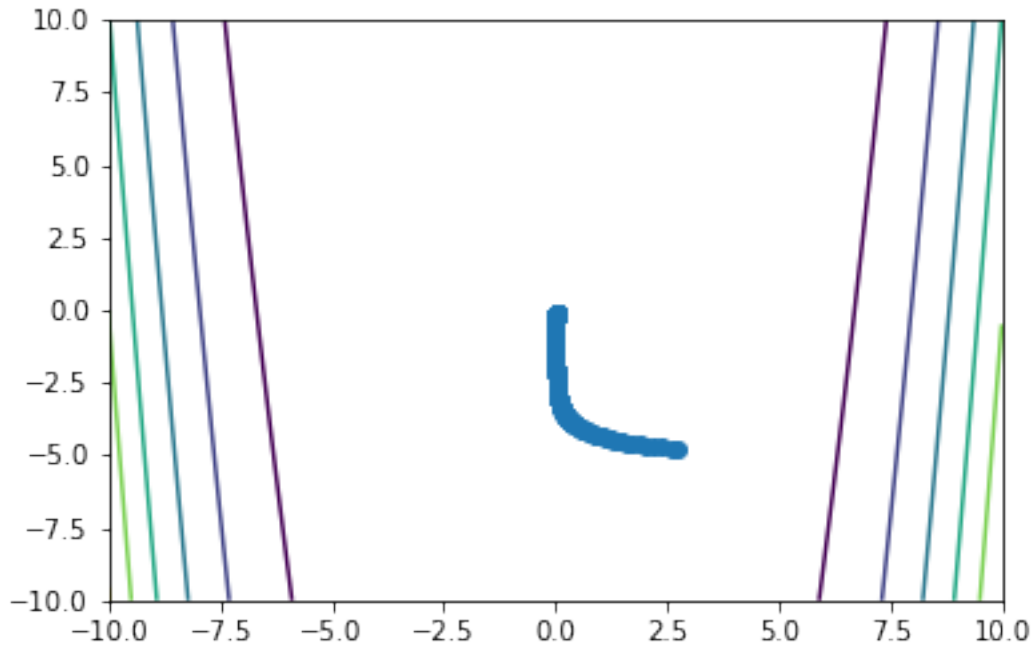
CS = ax.contour(X, Y, Z)
x= traj[:,0]
y= traj[:,1]
plt.plot(x,y, '-o')

```

Initial point is [ 2.71320643 -4.79248051]

After iterations, the point is [ 0.01632369 -0.08119181]

Out[119]: [<matplotlib.lines.Line2D at 0x13f27f630>]



In the plot above, I changed epsilon 10 times smaller. Initial points were the same, but the second one moved slower than the first one.

## 1 Problem 3

### 1.1 Part 1

Let  $x_i$  be  $i^{th}$  input for  $x_1, x_2, \dots, x_n$  and  $x_i$  is a 2D vector ( $2 \times 1$ ),  $W_1$  be the first weight matrix

$$W_1 = \begin{bmatrix} w_{111} & w_{121} \\ w_{112} & w_{122} \\ w_{113} & w_{123} \end{bmatrix}$$

$W_2$  be the second weight matrix

$$W_2 = \begin{bmatrix} w_{211} & w_{221} & w_{231} \\ w_{212} & w_{222} & w_{232} \end{bmatrix}$$

$W_3$  be the third weight matrix

$$W_3 = \begin{bmatrix} w_{311} & w_{321} \\ w_{312} & w_{322} \\ w_{313} & w_{323} \end{bmatrix}$$

Then, with loss function

$$z_1 = w_1 x + b_1, h_1 = f(z_1)$$

$$z_2 = w_2 h_1 + b_2, h_2 = f(z_2)$$

$$a = w_3 h_2 + b_3, \hat{y} = g(a)$$

$$L(y, \hat{y}) = y_1 \log(\hat{y}_1) + y_2 \log(\hat{y}_2) + y_3 \log(\hat{y}_3)$$

where  $f$  and  $g$  are Relu and softmax function respectively. Take gradient with respect to each parameters, we have

$$\frac{\delta L}{\delta a} = \frac{\delta L}{\delta \hat{y}} \frac{\delta \hat{y}}{\delta a} = \hat{y} - y$$

$$\frac{\delta L}{\delta b_3} = \frac{\delta L}{\delta a} \frac{\delta a}{\delta b_3} = (\hat{y} - y)$$

$$\frac{\delta L}{\delta w_3} = \frac{\delta L}{\delta a} \frac{\delta a}{\delta w_3} = (\hat{y} - y) h_2$$

$$\frac{\delta L}{\delta b_2} = \frac{\delta L}{\delta a} \frac{\delta a}{\delta h_2} \frac{\delta h_2}{\delta z_2} \frac{\delta z_2}{\delta b_2} = (\hat{y} - y) * w_3 * \mathbb{1}(z_2 > 0) * 1$$

$$\frac{\delta L}{\delta w_2} = \frac{\delta L}{\delta a} \frac{\delta a}{\delta h_2} \frac{\delta h_2}{\delta z_2} \frac{\delta z_2}{\delta w_2} = (\hat{y} - y) * w_3 * \mathbb{1}(z_2 > 0) * h_1$$

$$\frac{\delta L}{\delta b_1} = \frac{\delta L}{\delta a} \frac{\delta a}{\delta h_2} \frac{\delta h_2}{\delta z_2} \frac{\delta z_2}{\delta h_1} \frac{\delta h_1}{\delta z_1} \frac{\delta z_1}{\delta b_1} = (\hat{y} - y) * w_3 * \mathbb{1}(z_2 > 0) * w_2 * \mathbb{1}(z_1 > 0) * 1$$

$$\frac{\delta L}{\delta w_1} = \frac{\delta L}{\delta a} \frac{\delta a}{\delta h_2} \frac{\delta h_2}{\delta z_2} \frac{\delta z_2}{\delta h_1} \frac{\delta h_1}{\delta z_1} \frac{\delta z_1}{\delta w_1} = (\hat{y} - y) * w_3 * \mathbb{1}(z_2 > 0) * w_2 * \mathbb{1}(z_1 > 0) * x$$

## 1.2 Part 2

In [377]: `def grad_f(X,H1,H2,Y,Y_hat,parameters):`

`# Unpacking parameters`

`W1,b1,W2,b2,W3,b3 = parameters`

`# Gradients - ReLU`

`dW3 = np.dot((Y_hat - Y).T,H2)`

`db3 = np.reshape((Y_hat - Y).sum(axis=0),(3,1))`

`db2 = np.reshape((np.dot((Y_hat - Y),W3)* (H2 > 0)).sum(axis = 0),(2,1))`

`dW2 = np.dot((np.dot((Y_hat - Y),W3)* (H2 > 0)).T,H1)`

`db1 = np.reshape(np.sum(np.dot(np.dot((Y_hat - Y),W3)* (H2 > 0),W2)*(H1>0),axis = 0),(1,3))`

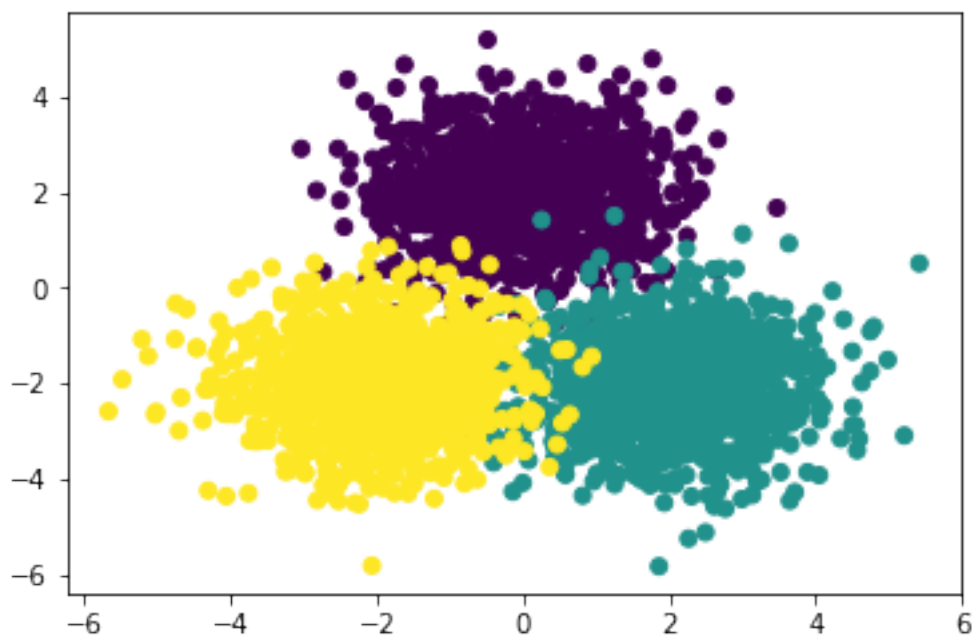
`dW1 = np.dot((np.dot(np.dot((Y_hat - Y),W3)* (H2 > 0),W2)*(H1>0)).T,X)`

`return [dW1,db1,dW2,db2,dW3,db3]`

### 1.3 Part 3

```
In [378]: covm = [[1, 0], [0, 1]]
x = np.concatenate((np.random.multivariate_normal(mean = [0, 2], cov=covm, size = 1000),
                    np.random.multivariate_normal(mean = [2,-2], cov=covm, size =1000),
                    np.random.multivariate_normal(mean = [-2, -2], cov=covm, size =1000)))
y = np.zeros(3000)
y[1000:2000] = 1
y[2000:3000] = 2
plt.scatter(x[:,0], x[:,1], c=y)
```

Out[378]: <matplotlib.collections.PathCollection at 0x118d3fc18>



### 1.4 Part 4

```
In [379]: Y = pd.Series(y)
Y = pd.get_dummies(Y).as_matrix()
X = x
```

/Users/FrankWang/anaconda3/lib/python3.6/site-packages/ipykernel\_launcher.py:2: FutureWarning:

```
In [380]: W1 = np.array([[1,0],[-1,0],[0,.5]])
W2 = np.array([[1,0,0],[-1,-1,0]])
W3 = np.array([[1,1],[0,0],[-1,-1]])
```

```

b1 = np.array([[0],[0],[1]])
b2 = np.array([[1],[-1]])
b3 = np.array([[1],[0],[0]])

In [381]: ReLu = np.vectorize(lambda z: np.fmax(0,z))
softmax = lambda z: np.exp(z)/(np.sum(np.exp(z),axis=1))[:,np.newaxis]
def predict(Y_hat):
    return np.argmax(Y_hat, axis=1)

In [382]: def loss(y, y_hat):
    tot = y * np.log(y_hat)
    return -tot.sum()
def forward(X,parameters):
    # Unpacking parameters
    W1,b1,W2,b2,W3,b3 = parameters
    # Forward pass
    a1 = np.dot(x,W1.T) +b1.T
    H1 = ReLu(a1)
    a2 = np.dot(H1,W2.T) + b2.T
    H2 = ReLu(a2)
    a3 = np.dot(H2,W3.T) + b3.T
    Y_hat = softmax(a3)
    return H1,H2,Y_hat

In [383]: def grad_descent(x, y, parameters, iterations=10, learning_rate=1e-2):
    point = parameters
    trajectory = [point]
    losses = [loss(y, forward(X,point)[2])]

    for i in range(iterations):
        H1,H2,Y_hat = forward(x,point)
        grad = grad_f(x,H1,H2, y,Y_hat,point)
        point = np.subtract(point, [i*learning_rate for i in grad])
        trajectory.append(point)
        losses.append(loss(y,forward(X,point)[2]))
    return (np.array(trajectory), losses)

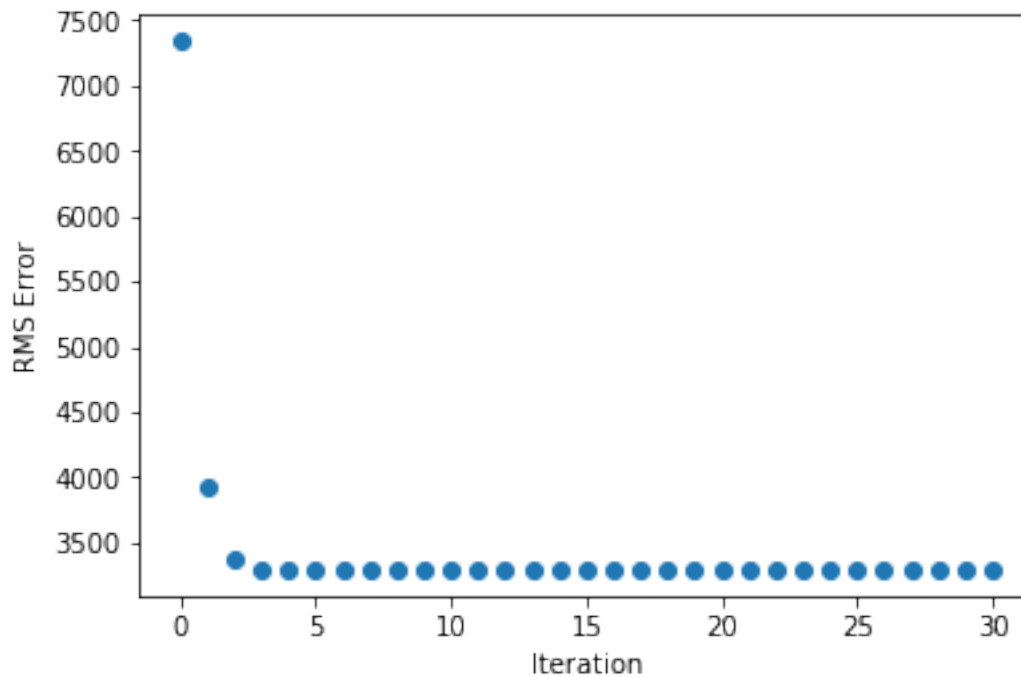
In [384]: parameters = [W1,b1,W2,b2,W3,b3]

In [385]: traj, losses = grad_descent(X,Y,parameters,iterations=30,learning_rate=1e-3)

In [386]: plt.plot(losses,'o')
plt.ylabel("RMS Error")
plt.xlabel("Iteration")

Out[386]: Text(0.5,0,'Iteration')

```



## 1.5 Part 5

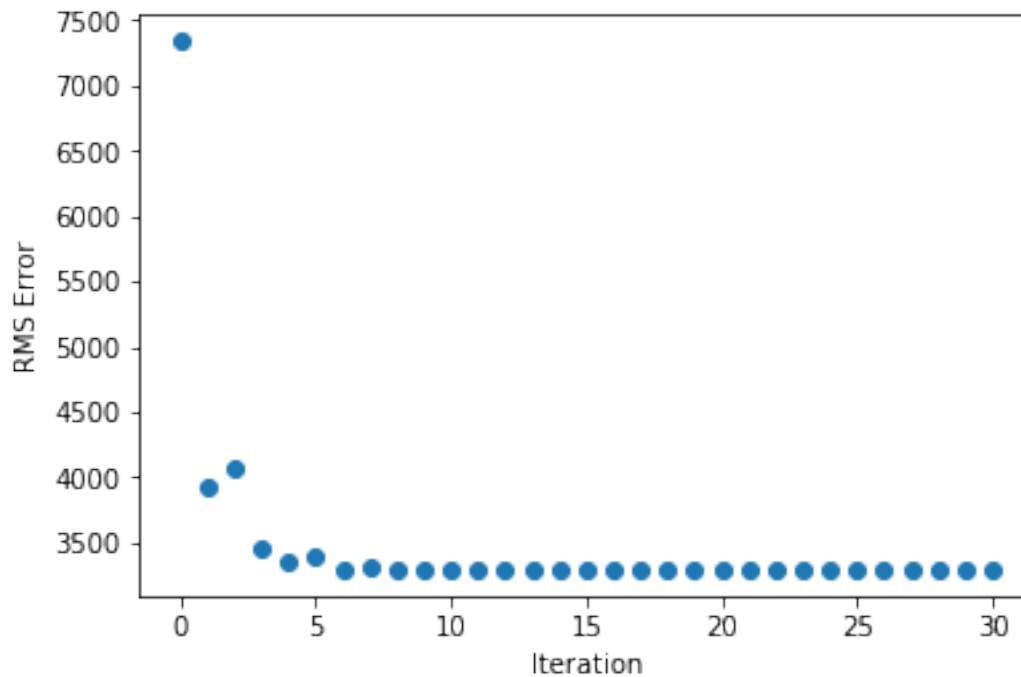
```
In [387]: def grad_descent_with_momentum(x, y, parameters, iterations=10, alpha=.9, epsilon=10):
    point = parameters
    trajectory = [point]
    losses = [loss(y, forward(X,point)[2])]
    v = [np.zeros(np.shape(i)) for i in parameters]

    for i in range(iterations):
        H1,H2,Y_hat = forward(x,point)
        grad = grad_f(x,H1,H2, y,Y_hat,point)
        #v = alpha*v + epsilon*grad
        v = np.add([alpha*i for i in v], [i*epsilon for i in grad])
        point = point - v
        trajectory.append(point)
        losses.append(loss(y,forward(X,point)[2]))
    return (np.array(trajectory), losses)

In [388]: traj, losses = grad_descent_with_momentum(X,Y,parameters,iterations=30,alpha=.5,epsilon=10)

In [389]: plt.plot(losses,'o')
           plt.ylabel("RMS Error")
           plt.xlabel("Iteration")

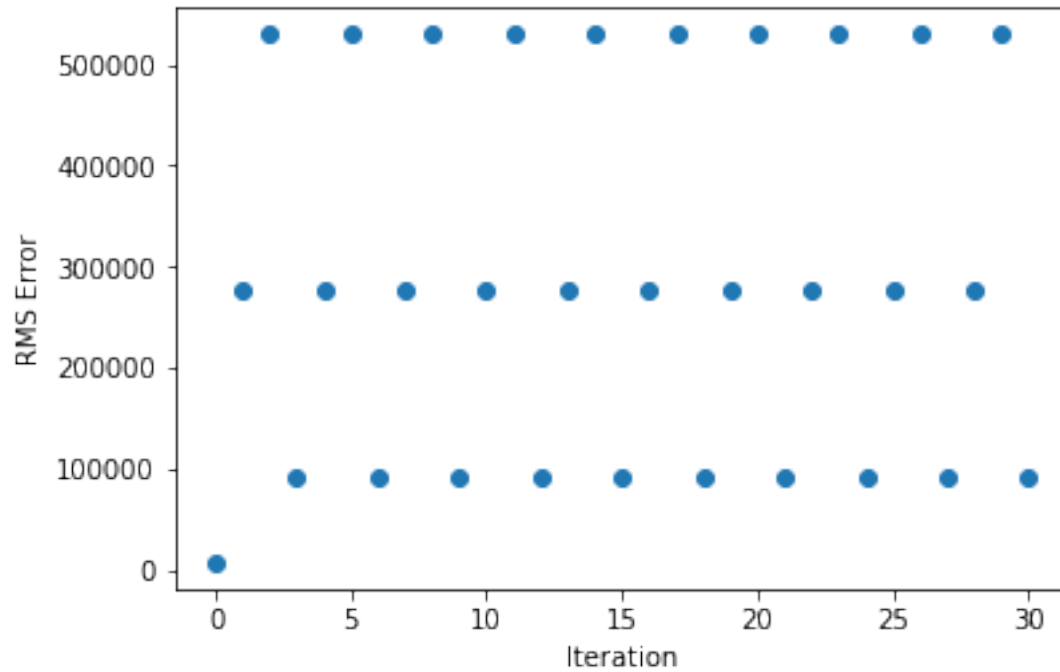
Out[389]: Text(0.5,0,'Iteration')
```



Both gradient descent and Momentum gradient descent have similar rate of convergence. But tuning parameters are important because smaller tuning parameter (learning rate in gradient descent and epsilon in Momentum gradient descent) will have a smaller step at each iteration and have a slower rate of convergence. However, if the tuning parameter is very large, it will skip the optimal point and become not stable. For example in the following plot, I change learning rate to 0.1 in gradient descent.

```
In [392]: traj, losses = grad_descent(X,Y,parameters,iterations=30,learning_rate=0.1)
          plt.plot(losses,'o')
          plt.ylabel("RMS Error")
          plt.xlabel("Iteration")
```

```
Out[392]: Text(0.5,0,'Iteration')
```



In the plot above, because learning rate is too large, it skipped the optimal point at each iteration and error is jumping around at each iteration. Hence, it is very important to choose a good tuning parameter